



TITANIC



Case study



Agenda

1. Dataset analysis
2. Preprocessing techniques
3. Output features, correlations and graphs
4. KNN classifier
5. Random forest classifier
6. Decision tree classifier

On April 15, 1912, after colliding with an iceberg, “unsinkable” Titanic sank resulting in death of 1502 out of 2224 passengers and crew.”

Dataset analysis

The dataset contains data of 891 passengers:


- PassengerId, Name, Sex, Age
- PClass - Class of the passenger (1, 2 or 3)
- Ticket - The ticket number (eg. “PC 17955”)
- Fare - How much did the passenger pay (eg. “14.45”)
- Cabin - Contains the cabin number (eg. “C85”)
- Embarked - Port of embarkation (S, C or Q)
- Parch - number of parents and children
- SibSp - number of siblings and spouses
- Survived - whether the passenger survived (1 or 0)

Total: 11 features and 1 binary decision label

Raw correlations

- Pclass and Fare
- SibSp and Parch
- Pclass and Age
- Survived and PClass
- Age and SibSp

- Better class = higher fare
- More parents = more siblings
- Better class = higher age
- Better class = survived
- Lower age = more siblings



Preprocessing Techniques

- a) Standardization (Fare)
- b) Handling missing values (Age, Embarked)
- c) One-hot encoding (Embarked, Title, Sex)
- d) Discretization (AgeBin)
- e) Derived attributes (FamilySize, Title)
- f) Dropped columns (Name, Age, SibSp, Parch)

| | Pclass | Fare | AgeBin | FamilySize | Embarked_Q | Embarked_S | Title_Miss | Title_Mr | Title_Mrs | Title_Other | Sex_male |
|---|--------|-----------|--------|------------|------------|------------|------------|----------|-----------|-------------|----------|
| 0 | 3 | -0.531122 | 2 | 2 | False | True | False | True | False | False | True |
| 1 | 1 | 1.099279 | 2 | 2 | False | False | False | False | True | False | False |
| 2 | 3 | -0.513935 | 2 | 1 | False | True | True | False | False | False | False |
| 3 | 1 | 0.636300 | 2 | 2 | False | True | False | False | True | False | False |
| 4 | 3 | -0.510753 | 2 | 1 | False | True | False | True | False | False | True |

Output features - explanation


Output feature analysis:

- PClass - unchanged
- Fare - standardized *Fare* column
- AgeBin - discretized *Age* column
- FamilySize - Size of family (sum of SibSp and ParCh columns)
- Embarked_Q, Embarked_S - port of embarkment expressed as hot label coding
- Title_Miss, Title_Mrs, Title_Mr, Title_Other - one hot encoded title
- Sex_male - sex expressed as one hot encoding
- Survived - unchanged

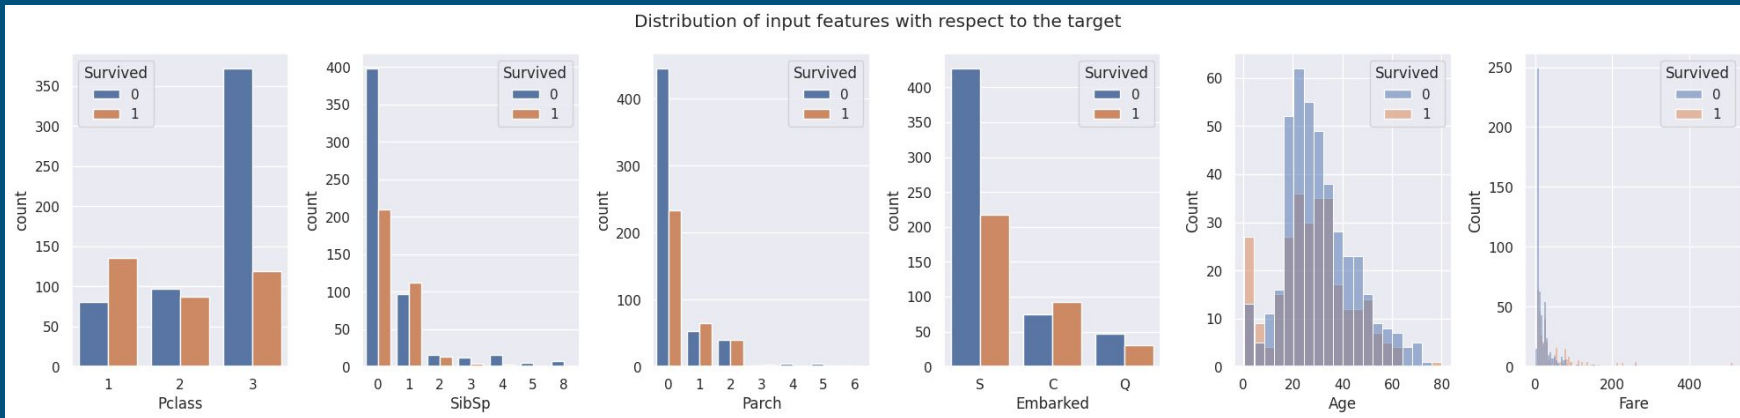
Preprocessing result - no missing values

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Pclass          891 non-null   int64
1   Fare            891 non-null   float64
2   AgeBin          891 non-null   category
3   FamilySize      891 non-null   int64
4   Embarked_Q      891 non-null   bool
5   Embarked_S      891 non-null   bool
6   Title_Miss      891 non-null   bool
7   Title_Mr        891 non-null   bool
8   Title_Mrs       891 non-null   bool
9   Title_Other     891 non-null   bool
10  Sex_male        891 non-null   bool
dtypes: bool(7), category(1), float64(1), int64(2)
memory usage: 28.1 KB
```

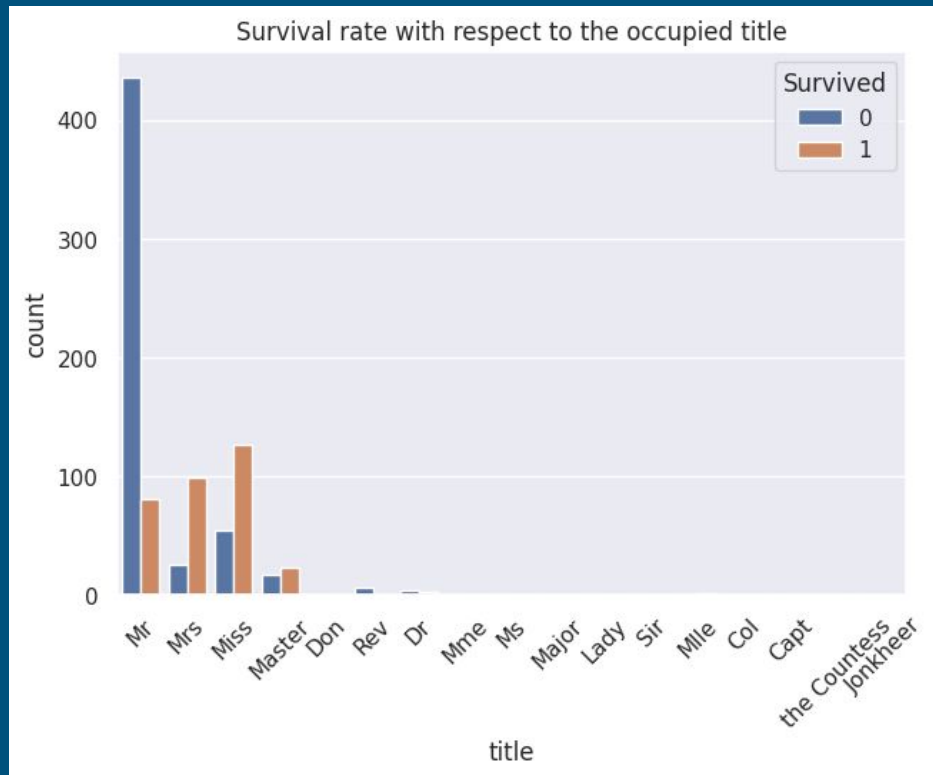
Distribution of target variable




Distribution of input features with respect to target



Dependency between titles and survival status



Survival Rate with respect to FamilySize



K-Nearest Neighbours (KNN)

```
k_values = range(1, 21)
metric_values = ['euclidean', 'manhattan', 'chebyshev']
search_method_values = ['ball_tree', 'kd_tree', 'brute']
```

Overall:

- Manhattan gives best results
- $K \sim 10$
- No need for complicated search methods
- 84% best accuracy (average of 10 splits 20%)

| k | metric | search_method | accuracy |
|----|-----------|---------------|----------|
| 10 | manhattan | brute | 0.840782 |
| 12 | manhattan | brute | 0.839665 |
| 11 | manhattan | ball_tree | 0.839106 |
| 7 | manhattan | ball_tree | 0.839106 |
| 7 | manhattan | brute | 0.837989 |
| 9 | manhattan | brute | 0.835754 |
| 20 | euclidean | ball_tree | 0.834078 |
| 20 | manhattan | ball_tree | 0.833520 |
| 14 | manhattan | brute | 0.831844 |
| 10 | euclidean | brute | 0.830726 |

K-Nearest Neighbours (KNN) - more parameters

- K is higher ~19
For top 3
- Again manhattan
Is the best
- Tree search methods
more common
- Higher accuracy 85.2%
(average of 10 splits 20%)
- 86.1% for 10% split

```
k_values = range(1, 21)
metric_values = ['chebyshev', 'minkowski']
p_values = [1, 2, 3, 4, 5] # 1 is manhattan, 2 is euclidean
search_method_values = ['ball_tree', 'kd_tree', 'brute']
weights_values = ['uniform', 'distance']
leaf_size_values = [10, 20, 30, 40, 50] # for trees
```

| k | metric | search_method | weight | leaf_size | p | accuracy |
|----|-----------|---------------|----------|-----------|---|-----------|
| 18 | minkowski | kd_tree | uniform | 50 | 1 | 85.195531 |
| 20 | minkowski | ball_tree | uniform | 20 | 1 | 85.139665 |
| 18 | minkowski | kd_tree | distance | 10 | 1 | 84.972067 |
| 8 | minkowski | kd_tree | uniform | 20 | 1 | 84.916201 |
| 10 | minkowski | ball_tree | uniform | 20 | 1 | 84.860335 |
| 14 | minkowski | kd_tree | uniform | 40 | 1 | 84.804469 |
| 10 | minkowski | ball_tree | uniform | 40 | 2 | 84.692737 |
| 10 | minkowski | kd_tree | uniform | 20 | 1 | 84.692737 |
| k | metric | search_method | weight | leaf_size | p | accuracy |
| 12 | minkowski | kd_tree | uniform | 50 | 1 | 86.111111 |
| 9 | minkowski | kd_tree | uniform | 40 | 1 | 85.555556 |
| 13 | minkowski | brute | uniform | 10 | 1 | 85.444444 |

Random Forests!

RandomForestClassifier

```
class sklearn.ensemble.RandomForestClassifier(n_estimators=100, *,
criterion='gini', max_depth=None, min_samples_split=2, min_samples_leaf=1,
min_weight_fraction_leaf=0.0, max_features='sqrt', max_leaf_nodes=None,
min_impurity_decrease=0.0, bootstrap=True, oob_score=False, n_jobs=None,
random_state=None, verbose=0, warm_start=False, class_weight=None,
ccp_alpha=0.0, max_samples=None, monotonic_cst=None) #
```

[\[source\]](#)

Random Forest - hyperparameter tuning

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

grid_search = GridSearchCV(estimator=RandomForestClassifier(random_state=42),
                           param_grid=param_grid,
                           cv=3,
                           verbose=3,
                           n_jobs=-1)

grid_search.fit(X_train, y_train)
```

Random Forest

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

params = {
    'bootstrap': True,
    'max_depth': 10,
    'min_samples_leaf': 1,
    'min_samples_split': 10,
    'n_estimators': 100
}

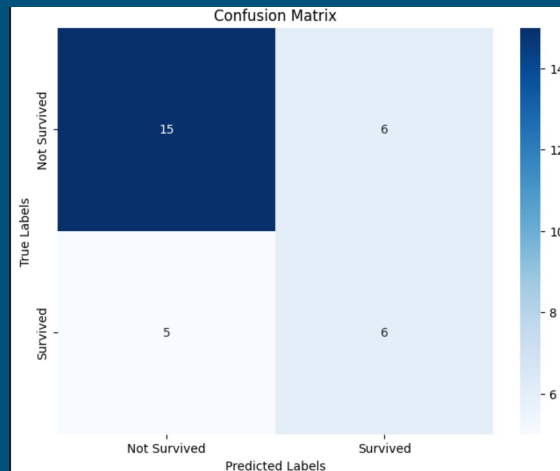
rf_classifier = RandomForestClassifier(**params)
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the Random Forest model: {accuracy:.2f}")
```

✓ 0.9s

Accuracy of the Random Forest model: 0.69

- Simple
- Sad accuracy :(



Decision Tree - sklearn

```
class sklearn.tree.DecisionTreeClassifier(*, criterion='gini', splitter='best',  
max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0,  
max_features=None, random_state=None, max_leaf_nodes=None, min_impurity_decrease=0.0,  
class_weight=None, ccp_alpha=0.0, monotonic_cst=None) \[source\]
```

We will be interested in:

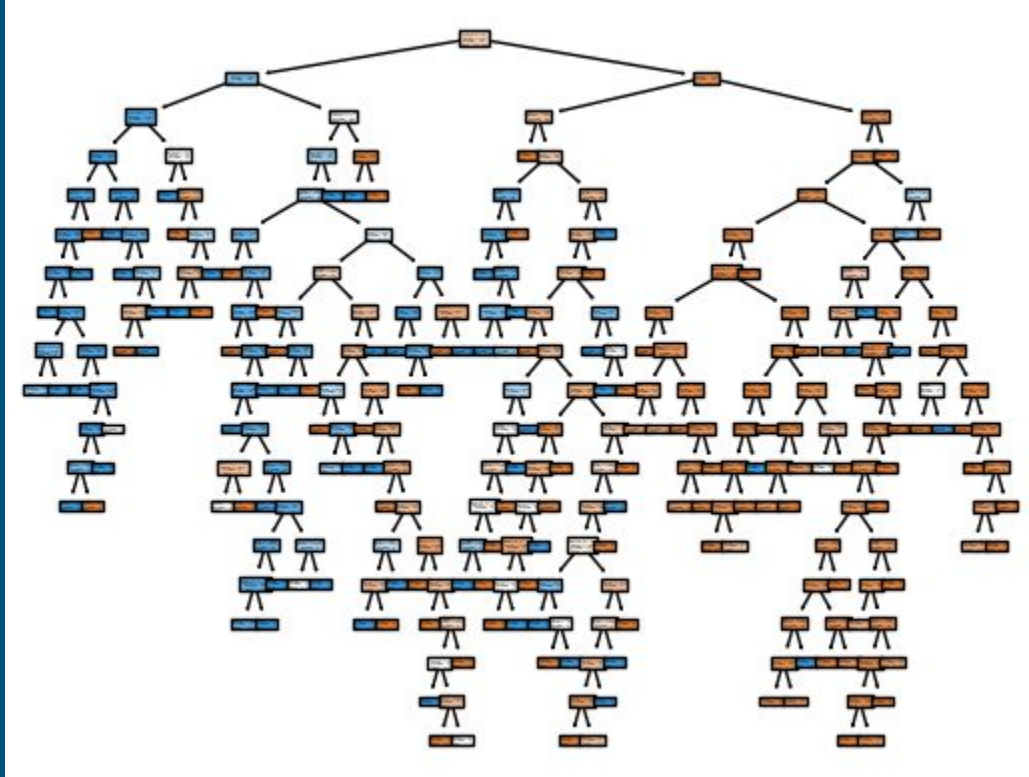
- | | |
|--------------|----------------------|
| a) criterion | d) min_samples_split |
| b) splitter | e) min_samples_leaf |
| c) max_depth | f) ccp_alpha |

Decision tree - how to use it?

```
clf = tree.DecisionTreeClassifier(criterion = "entropy")
clf = clf.fit(X_train, y_train)
tree.plot_tree(clf, filled = True, feature_names=X_train.columns)

y_tree = clf.predict(X_test)
accuracy = accuracy_score(y_tree, y_test)
print(f"Accuracy obtained for decision tree: {round(100*accuracy, 2)}%")
```

Decision Tree - First case (default)

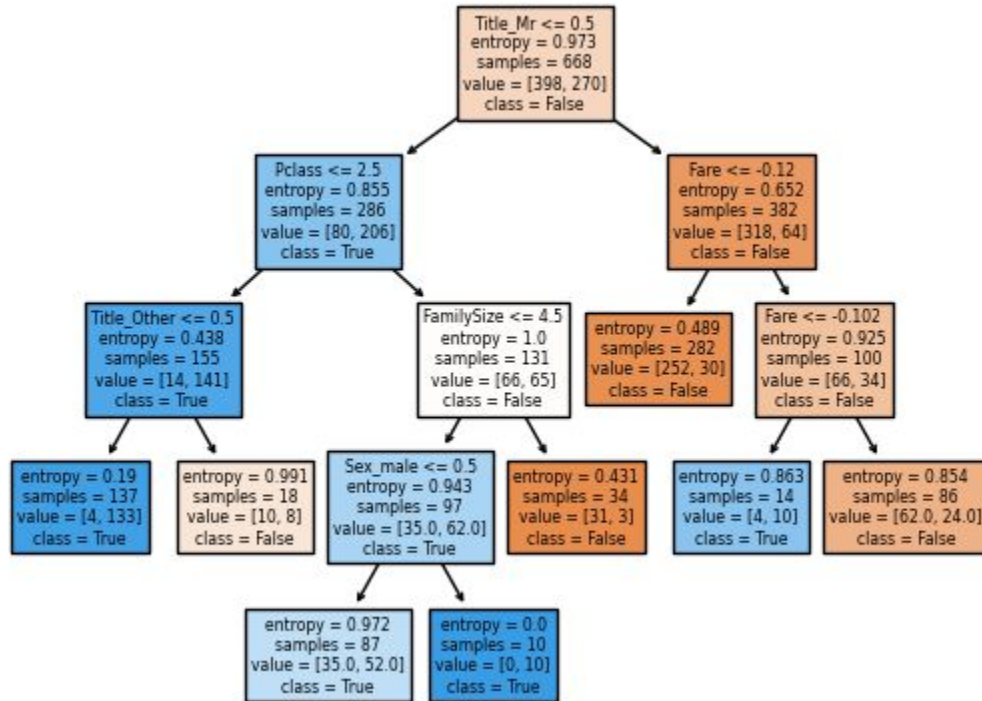


Parameters:

- a) Criterion = "entropy"
- b) Other parameters = default

Accuracy obtained (avg): 78.54%

Decision Tree - Second case (custom)




Parameters:

- a) criterion: "entropy"
- b) min_samples_split = 25
- c) min_samples_leaf = 5
- d) ccp_alpha = 0.01
- e) Max_depth = 5

Accuracy obtained (avg): 81.42%

Decision Tree - Third case (trained, no ccp)

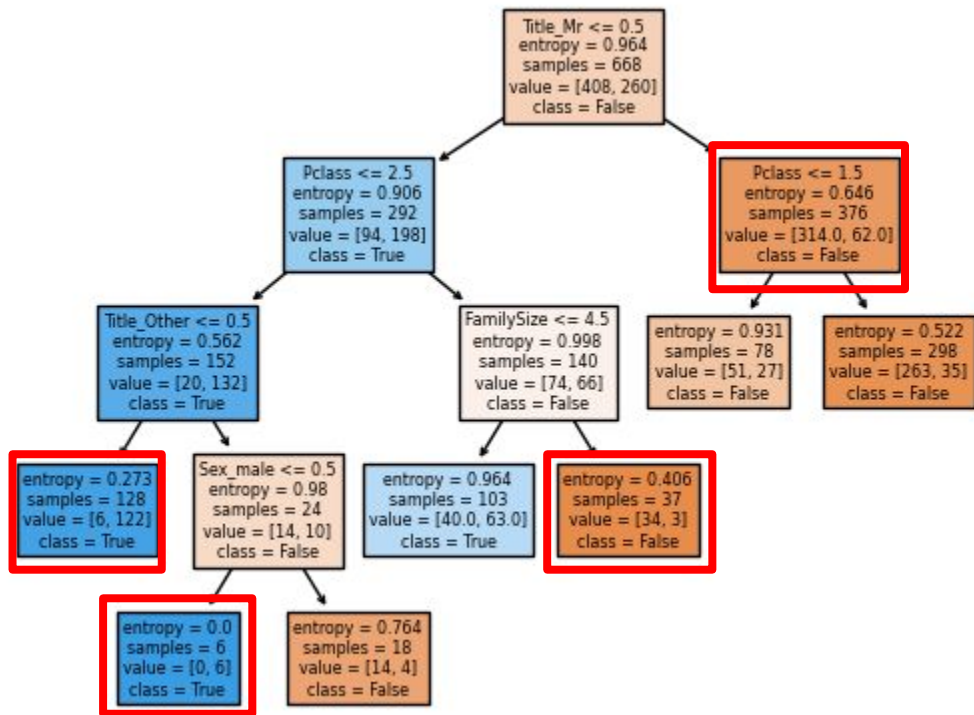


Parameters:

- a) criterion: "entropy"
- b) min_samples_split = 2
- c) min_samples_leaf = 2
- d) ccp_alpha = 0
- e) Max_depth = 11

Accuracy obtained: 87.0%

Decision Tree - Fourth case (trained, ccp)

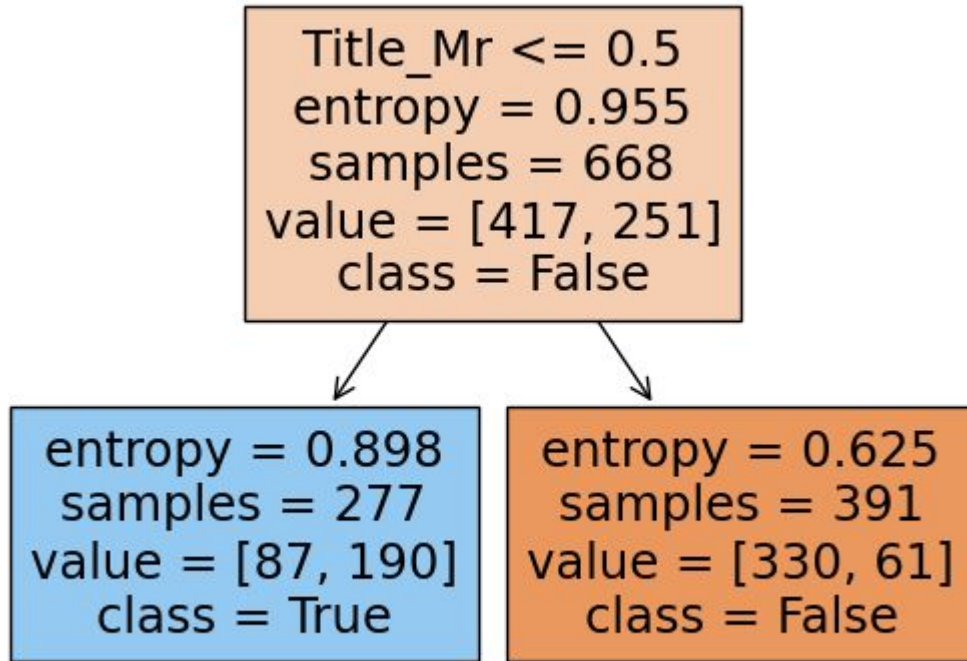


Parameters:

- a) criterion: "entropy"
- b) min_samples_split = 5
- c) min_samples_leaf = 2
- d) ccp_alpha = 0.01
- e) Max_depth = 4

Accuracy obtained: 85.65%

Be careful with ccp_alpha!



Parameters:

- a) criterion: "entropy"
- b) min_samples_split = 5
- c) min_samples_leaf = 2
- d) ccp_alpha = 0.1
- e) Max_depth = 4


Accuracy obtained: 77.28%

More stats...

For more statistics (accuracies, confusion matrices etc) and code implementation go to jupyter notebook!

Conclusions

- Model Performance
- Feature Insights
- Preprocessing Impact



Sources

- a) scikit-learn.org (documentation)
- b) wikipedia.org
- c) kaggle.com

Thank You

- Mateusz Idziejczak 155842
- Mateusz Stawicki 155900
- Kuba Czech 156035